

Artificial Neural Networks (ANNs)

1. Activation Functions in ANNs

Activation functions introduce nonlinearity into the network, enabling it to learn complex patterns. Common activation functions include:

- **Sigmoid**

- **Formula:** $\sigma(x) = \frac{1}{1+e^{-x}}$

- **Properties:** Squashes input to (0, 1); historically used in binary classification but can suffer from vanishing gradients.

- **Hyperbolic Tangent (tanh)**

- **Formula:** $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

- **Properties:** Output ranges from -1 to 1; zero-centered which often speeds up convergence compared to sigmoid.

- **Rectified Linear Unit (ReLU)**

- **Formula:** $f(x) = \max(0, x)$
 - **Properties:** Computationally efficient; helps mitigate the vanishing gradient problem but can lead to “dying ReLU” where neurons become inactive.

- **Leaky ReLU and Parametric ReLU (PReLU)**

- **Formula:** $\text{Leaky ReLU}(x) = \begin{cases} x, & \text{if } x \geq 0 \\ \alpha x, & \text{if } x < 0 \end{cases}$ (with a small constant α , e.g., 0.01)

- **Properties:** Addresses the dying ReLU issue by allowing a small, non-zero gradient when $x < 0$.

- **Exponential Linear Unit (ELU) and Scaled ELU (SELU)**

- **Properties:** Aim to bring mean activations closer to zero, which can speed up learning and improve robustness.

- **Softmax**

- **Properties:** Used in the output layer for multi-class classification, providing a probability distribution across classes.

- **Formula:** $\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}}$

2. Types of Layers in ANNs

Layers in ANNs are the building blocks that process inputs and pass them through the network. Common types include:

- **Input Layer:**
 - The first layer of the network that receives raw input data.
- **Fully Connected (Dense) Layer:**
 - Each neuron is connected to every neuron in the previous layer; used in many standard neural network architectures.
- **Hidden Layers:**
 - Layers between the input and output that perform complex transformations on the data. They can be fully connected or take other specialized forms.
- **Convolutional Layers:**
 - Used primarily in image and signal processing tasks; they apply convolution filters to capture spatial hierarchies.
- **Pooling Layers:**
 - Typically used after convolutional layers to reduce spatial dimensions (e.g., max pooling, average pooling).
- **Recurrent Layers:**
 - Employed in sequence processing tasks; examples include LSTM (Long Short-Term Memory) and GRU (Gated Recurrent Unit) layers.
- **Normalization Layers:**
 - Layers like Batch Normalization or Layer Normalization help stabilize and accelerate training.
- **Dropout Layers:**
 - Randomly drop neurons during training to prevent overfitting by ensuring the network does not rely too much on any one neuron.

3. Types of Optimizers

Optimizers adjust the weights of the network during training by minimizing the loss function. Popular optimizers include:

- **Stochastic Gradient Descent (SGD):**
 - **Description:** Updates weights using the gradient computed on a mini-batch of data.

- **Variants:** Often used with momentum, which adds a fraction of the previous update to the current one for smoother convergence.
 - **Momentum:**
 - **Description:** Accelerates SGD by navigating along relevant directions and dampening oscillations.
 - **AdaGrad:**
 - **Description:** Adapts the learning rate for each parameter based on historical gradients; performs well for sparse data.
 - **RMSProp:**
 - **Description:** Modifies AdaGrad by using a moving average of squared gradients to dampen the aggressive, monotonically decreasing learning rates.
 - **Adam (Adaptive Moment Estimation):**
 - **Description:** Combines the ideas of momentum and RMSProp, keeping an exponentially decaying average of past gradients and squared gradients.
 - **Variants:** Adamax, Nadam (which incorporates Nesterov momentum) are some examples.
-

4. Types of Loss Functions

Loss functions measure the error between the network's predictions and the actual target values. The choice depends on the task:

- **Mean Squared Error (MSE):**
 - **Usage:** Commonly used in regression tasks.
 - **Formula:**
$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2.$$
- **Mean Absolute Error (MAE):**
 - **Usage:** Also used for regression; less sensitive to outliers compared to MSE.
- **Binary Cross-Entropy (Log Loss):**
 - **Usage:** Used for binary classification problems.
 - **Formula:**
$$L = -\frac{1}{n} \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)].$$
- **Categorical Cross-Entropy:**
 - **Usage:** Used for multi-class classification problems.
- **Hinge Loss:**

- **Usage:** Common in training classifiers such as support vector machines (SVMs) or in certain neural network contexts.
 - **Huber Loss:**
 - **Usage:** Combines MSE and MAE, providing robustness to outliers while being differentiable at zero.
-